

Embedded Kafka

Thomas Ambrogini

December 5, 2023

Contents

1	Embedded Systems	1
1.1	Microcontoller, Microprocessor and SoC	2
1.1.1	Microcontroller	2
1.1.2	Microprocessor	3
1.1.3	SoC	3
1.2	Toolchain	3
1.2.1	Cross Compiler	4
1.2.2	Linker	5
1.2.3	Make	6
1.2.4	Debugger	6
1.3	Boot	7
1.3.1	Multi-stage Bootloader	7
1.3.2	Multi core bootloading	8
1.4	Operating Systems	9
1.4.1	Real Time Operating Systems	10
1.4.2	Linux Embedded Systems	11

List of Tables

List of Figures

1.1	Microprocessor and microcontroller architectures	2
1.2	Cross Compilation process	4
1.3	Multi stage bootloader	8
1.4	Task states	11
1.5	Embedded Linux Architecture	12
1.6	Architecture of Real Time Linux	14
1.7	Yocto Stack	15

Chapter 1

Embedded Systems

Designing and writing software for embedded systems poses a different set of challenges than traditional high-level software development. This chapter gives an overview of this challenges.

Embedded systems are computing devices performing specific, dedicated tasks with no direct or continued user interaction. [8] Due to the variety of market and technologies, these objects have different shapes and sizes, in general, they are small and resource constrained.

One of the characteristics for which embedded systems are used is the capacity of delivering real time processing. Real time processing refers to external events within precise and predictable timeframes. In manufacturing, robotics, and automotive production, timely responses to changing conditions can mean the difference between operational success and failure. Embedded systems excel at this, ensuring that processes are executed with minimal latency, guaranteeing the coordination and synchronization of various components.

Embedded systems' ability to deliver real time processing capabilities is fundamental in automation industries. However, the importance of real time processing extends beyond the automation industry. In the healthcare sector, embedded systems can provide the delivery of life-saving treatments.

Most embedded systems are resource constrained, limiting memory and processing power. These constraints are a direct result of the need for compact, energy-efficient and cost-effective solutions. These limitations impose signif-

icant design challenges, forcing developers to optimize their code and data storate in order to respect those limits.

1.1 Microcontoller, Microprocessor and SoC

Embedded systems can rely on different integrated circuits (ICs) for executing applications, ranging from simpler, self-contained architectures to more advanced ones, which requires the integration of other components in order to work.

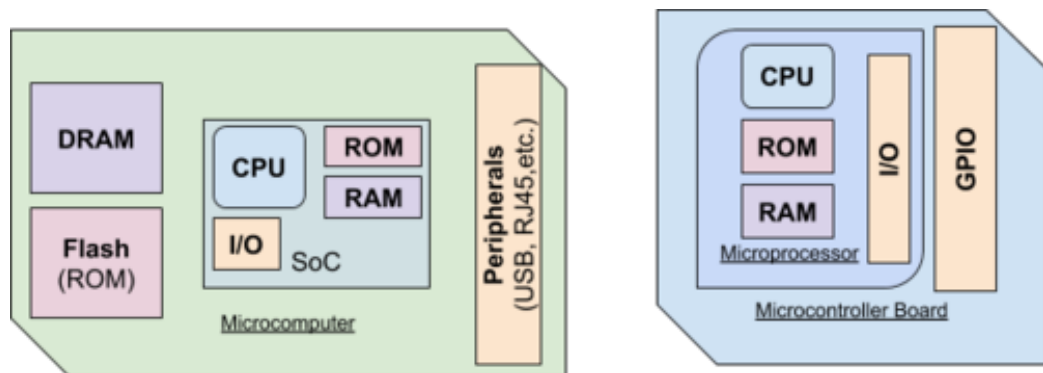


Figure 1.1: Microprocessor and microcontroller architectures

1.1.1 Microcontroller

Microcontrollers are compact, self-contained computing devices that find use in embedded systems. They are designed to perform specific tasks with great efficiency. A typical microcontroller incorporates a central processing unit (CPU), memory, input/output pins, and a range of peripherals within a single chip. This integration minimizes the need for external components, reducing cost, size, and power consumption. Microcontrollers excel in real-time control, making them suitable for applications where precise timing and responsiveness are crucial. Their versatility, energy efficiency, and cost-effectiveness have made them the preferred choice in an array of embedded systems, including robotics, automotive control systems, consumer electronics, and the Internet of Things (IoT).

1.1.2 Microprocessor

Microprocessors are the computational brains of modern computing systems. These general-purpose integrated circuits are designed to process data and execute instructions efficiently. Unlike microcontrollers, microprocessors do not integrate a wide array of peripherals, but they offer greater versatility in terms of the software and hardware components they can work with.

Microprocessors are commonly found in devices like personal computers, smartphones, and laptops, where their high processing power is a requirement. They excel in applications that demand data-intensive tasks, multitasking, and complex computation.

1.1.3 SoC

System-on-Chip (SoC) integrates multiple components into a single chip, just like microcontrollers. The main difference between SoCs and microcontrollers is that the latter can integrate also GPUs and various accelerators, making them more complex and expensive. The integration of these elements into a single chip enhances efficiency and reduces power consumption, making SoC really valuable in the embedded systems world, especially for the automation industry, which could require substantial processing power for certain applications.

They provide a heterogeneous architecture with different processing units into the single IC (CPU, GPU, DSP) which can have great performances in multitask and multithread applications.

1.2 Toolchain

Developers rely on a set of tools to compile, link, execute and debug software to a specific target. Building the firmware images for an embedded system relies on a similar set of tools, called a toolchain.

In order to create an executable for an embedded system, the process is more more complex than just creating the executable for a regular development environment. Since the architecture of the CPU is not the same of the

host machine, a compiler that generates machine code for the specific target architecture is needed. The process of creating executable for another architecture is called cross-compilation and the compiler needs some additional information to perform it.

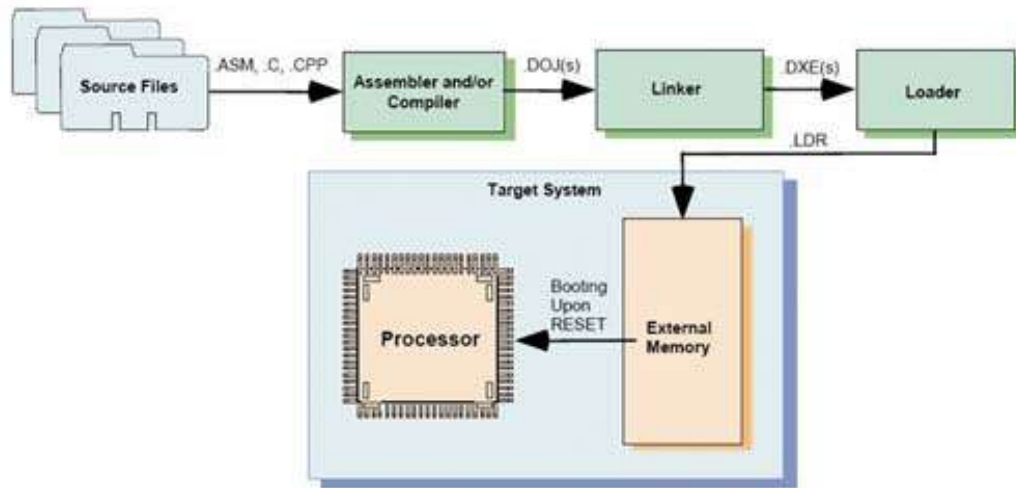


Figure 1.2: Cross Compilation process

1.2.1 Cross Compiler

The first element of the toolchain is the compiler, which is responsible for translating source code into machine code, which can be interpreted by a specific CPU. Each compiler can produce code for only one environment, since the source code is translated to machine specific instructions, which uses the registers and the address model of a specific architecture.

The compiler creates object files (.o) starting from the source code. These files contain compiled code for the target machine, which is still not executable, but can be linked to other object files to create the executable.

Object files contain instructions for the CPU and a symbol table, containing information about functions and variables of the program. These files have information about the functions and variables initial value.

1.2.2 Linker

The linker performs the creation of the executable file. It aggregates all the object files received as input created by the compiler and resolves all the meaning for every symbol and all the dependencies, finally producing the executable.

The standard executable format is called ELF (executable and linkable format) and it has been designed to represent programs on disk and other media. It can be executed by loading the information in RAM in specific addresses. ELF files are divided into sections, each one of them represents a different area of memory with information needed for the execution of the program. They also contain an header with a pointer to the different sections inside the file. The main sections are:

- `.text`: instructions of the program.
- `.rodata`: read only variables.
- `.data`: variables accessible in read and write.
- `.bss`: uninitialized variables which can be accessed in read or write mode.

Read-only information can be loaded directly from flash in an embedded system, while data that need to be modified has to be in RAM on modifiable areas of memory.

In regular development environments, much of the complexity of the linking step is hidden, but the linker is configured, by default, to rearrange the compiled symbols into sections which can be loaded into the process' virtual addresses by the operating system.

In an embedded system, the linking process becomes more complex as compiled symbols have to be placed in physical addresses specific to the system's architecture. In order to specify in which areas of memory the sections have to be placed, a custom linker file has to be created. This file defines the memory layout of the executable bare metal application. Additionally, in linker files, custom sections can be used if they are required by the target system.

1.2.3 Make

GNU Make is a tool which controls the generation of executable and other non-source files of a program from the program's source files. [5] Make gets its knowledge on how to build binary images from a file called the makefile, which lists how to create each of the non-source files.

The makefile operates on rules and targets. Rules are a series of commands which instruct make on how to produce a target, representing a non-source file.

One of the advantages of using a build automation tool is that some of the targets could have implicit dependencies from other intermediate components, that are automatically resolved at compile time. It also reduce the time required for the compilation by compiling targets only when an update to the dependencies has been done.

1.2.4 Debugger

One of the most powerful tools within a toolchain is the debugger. Debuggers allows developer to test the runtime functionalities of an application, ensuring it operates as intended or locating the source of errors.

In the host environment, debugging an application consists of launching the debugger tool with an ELF file as argument or attaching it to a running application. The standard tool for debugging is called GDB.

However, in embedded development, the situation is slightly different since the application runs on a different machine than the one used for debugging. To address this, a version of GDB to debug remote platforms has been developed. The remote debug session requires an intermediate component which can translate the GDB commands to CPU specific instructions.

The debug of remote embedded platforms often require a peripheral like JTAG.

1.3 Boot

A bootloader is a piece of software that starts as soon as you power on the system. Its primary functionality is to initiate subsequent software components such as: an operating system, a bare metal application or in some cases another bootloader. In the embedded world, bootloaders are intimately tied to the underlying architecture of the SoC. They are typically securely stored in a non-volatile, protected on-chip memory, ensuring their availability at startup.

Upon execution, the bootloader execute a series of tasks. It begins by performing hardware checks to verify the integrity and functionality of the SoC's components. Then, it takes charge of initializing the processor and configuring essential system-on-chip registers.

Bootloaders are also important in ensuring security. They often serve as the starting point for establishing a Hardware Root-of-trust, which forms the foundation for securing the entire system.

1.3.1 Multi-stage Bootloader

In advanced SoCs, it's not uncommon for the bootloader to load another bootloader, creating a multi-stage bootloading process. In this case the bootloaders serve different purposes.

The initial bootloader, called the first stage bootloader, is typically stored in a secure, read-only memory for enhanced security and to maintain a consistent, known state each time the system powers on. The first stage bootloader is intentionally kept simple and performs only the essential configuration tasks required to bring up the system.

The secondary bootloader, or second stage bootloader, takes on a more complex and configurable role to serve the specific needs of the application. The second stage bootloader can be updated more easily and frequently compared to the first stage bootloader, making it more adaptable to changes and enhancements as the system evolves or new features are added.

The multi-stage approach to bootloading exploit the first stage for security and configuration, while the second one actually initializes the system, pro-

viding modularity and flexibility for future updates.

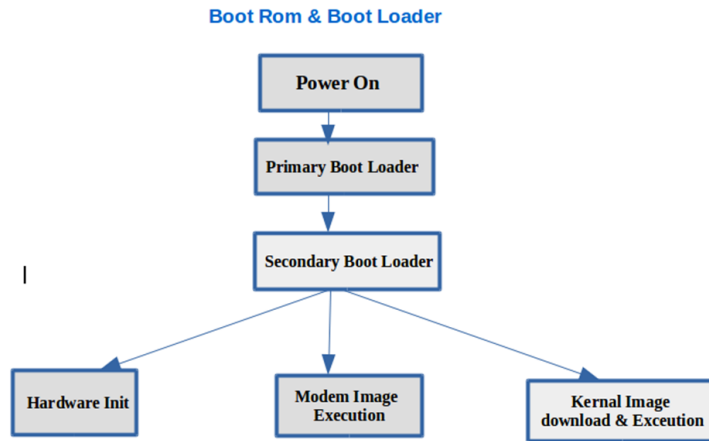


Figure 1.3: Multi stage bootloader

1.3.2 Multi core bootloading

Modern embedded systems have multi core architectures that must be initialized, and the process of bootloading must be carried out by multiple cores, working together.

Multi core bootloading can be seen as an extension of a multi-stage bootloader. In this scenario, the first stage bootloader may not even be aware of the presence of other cores. Its primary responsibility is to initiate the next-stage bootloader, which then takes charge of orchestrating the complex bootloading process across multiple cores.

The complexities which are introduced in this process are:

1. Image preparation: each processor core may have its own code and data segments. The images for each core need to be prepared in a specific format that accounts for their individual requirements.
2. Image format: the images for every core could be concatenated into a single image for the second stage bootloader to load. The SBL must be aware of the format used, if only one image is used, to correctly parse it.

3. Coordination: coordinating the bootloading process across multiple cores involves setting up the execution context for each core, initializing hardware components, and synchronizing the start of execution. The SBL takes on the role of managing these tasks to ensure a smooth and reliable multi-core startup.

1.4 Operating Systems

In the embedded world there are different kind of operating systems, not only realized by different companies, but also with different needs as objective. Embedded devices, usually, needs to satisfy real time requirements, which makes hard using a general operating system, since they can not provide this feature.

The essence of real time is not that a computer has to respond fast, but that it has to respond reliably fast. The requirement of real time programming is being able to quickly handle asynchronous events. [7]

For many years, the only solution used for these devices was the bare metal approach. In this case there isn't a real operating system which takes care of the management of fundamental operations. A bare metal application is considered to run directly on the hardware without respecting an external programming interface (e.g., the one given by an operating system), having direct access to CPU or microcontroller registers and without the security mechanisms of an OS.

Bare metal programming has the advantage of providing to the developers the highest grade of freedom, understanding exactly what every action will end up doing. Bare metal applications has the greatest possible degree of determinism and the resource consumption is optimized for the specific case. On the other hand, it becomes hard to manage large project with different tasks and multiple operations to perform.

Other than the bare metal approach, there is the possibility of running a Real Time Operating System, which provides support for multiple tasks and device drivers between the hardware and the application, stack for the network and security.

1.4.1 Real Time Operating Systems

As the complexity of tasks continue to increase for embedded devices the necessity for a RTOS has increased and always more embedded devices go towards this solution, leaving behind the bare metal approach. These systems gives an efficient solution to meet hard real time requirements, particularly in safety critical applications that requires the management of different safety applications on the same device.

The scheduler in a RTOS, which has the job of deciding which thread have to execute on the core in any moment, is made in such a way to guarantee deterministic execution pattern. Real Time schedulers achieve determinism by assigning a priority to each task. In this way the scheduler can always run the task with the highest priority. [4]

FreeRTOS

FreeRTOS is a real-time operating system kernel for embedded devices. It is designed to be small and simple, built with an emphasis on reliability and ease of use. It is ideally suited for real-time applications, including a mix of both soft and hard real time requirements.

FreeRTOS allows applications to be organized as a collection of independent threads of execution. On a single core processor, only one thread can execute in any given time. The kernel decides which thread should be executing by examining the priority assigned to each thread by the application designer. Typically, threads with hard real-time requirements have higher priorities, ensuring their execution ahead of threads with soft real-time requirements.

In a single core system, the CPU must divide the tasks into time slices so that they can appear to run concurrently. The scheduler in an operating system is charged with figuring out which task to run each time slice. In FreeRTOS, the default time slice is 1 millisecond, referred to as "tick". A hardware timer is configured to create an interrupt every 1 ms and the handler for that interrupt runs the scheduler, which chooses the task to run next. At each tick interrupt, the task with the highest priority is chosen to run. If the highest priority tasks have the same priority, they are executed in a round-robin fashion. If a task with an higher priority than the currently running task becomes available, then it will immediately run, without waiting for the

next tick.

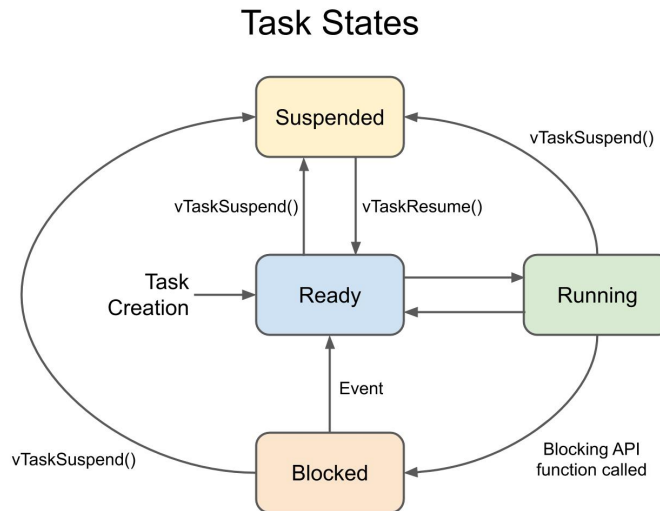


Figure 1.4: Task states

As soon as a task is created, it enters the Ready state, signaling their readiness to the scheduler for potential execution.

In multi core systems, the scheduler can schedule tasks based on the number of available cores in the system.

In addition, FreeRTOS provides features to simplify the embedded software development. It includes services for task scheduling, synchronization primitives, resource management, and interrupt handling.

1.4.2 Linux Embedded Systems

Embedded Linux is built on the same Linux kernel of every other linux systems. Other than the Linux kernel, embedded applications need additional packages, which depends on the distribution and can be chosen based on the necessities of the developer. [3]

Although most of the devices are small in size, there are some of them that can run a version of the GNU/Linux operating system. To have this possibility the device must of a reasonable amount of RAM and the hardware

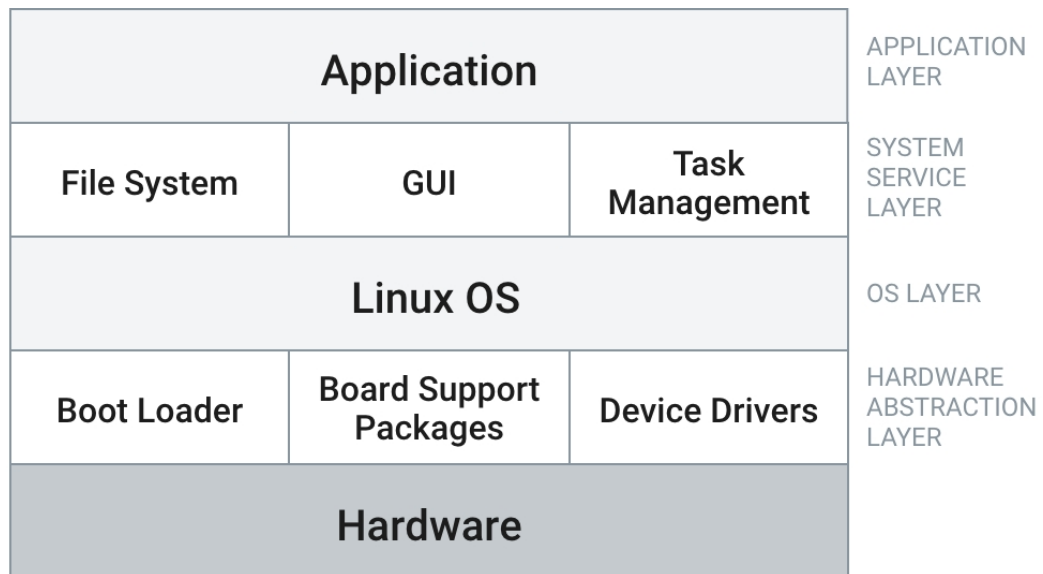


Figure 1.5: Embedded Linux Architecture

must support MMU, Memory Management Unit, to provide different virtual address spaces for every process.

The main advantage of using a system with Linux is that the opportunity for a tailored solution will be less, since the requirements, resource wise, for running Linux makes a system overkill for most of the applications. In this way, the software design can be simpler and it is possible to use existing solutions to embedded problems.

On the other hand, embedded devices have in many cases hard real time requirements, meaning that a series of operations must be accomplished in a short, measurable, predictable amount of time. To accomplish this Real-Time Operating Systems are used and it is almost impossible to achieve real time processing with embedded Linux as operating system, even if some patches to the kernel's scheduler have been applied to help meet these requirements.

Other application domains for embedded devices are low power consumptions, which could run on a battery or energy harvesting device. In this case, having an operating system increase the energy need of the device.

A reason for which linux could be a perfect candidate as an operating system for embedded devices is its modularity. Embedded developers have the possibility to customize their linux distribution, including only the necessary

parts for their use case. [2]

Linux Embedded Real Time

One way to improve real time performances in an operating system is to add extra preemption points, where the OS can switch to critical operations. However, this process worsen the overall performances of the operating system in the general case, which is what general systems are optimized for, not taking into consideration the worst-case scenario, making the system non-deterministic.

The solution to the problem is to decouple the real time part of the system from the general purpose kernel. It is possible to optimize the real time OS to meet deadlines, having the best performance of general computing. [1]

Different projects to integrate real time into linux have been realized (e.g., RTLinux, RTAI). Although, they are maintained by different people, most of the functionalities are the same between different projects:

- A small real time core.
- One shot and periodic timer support.
- Real time scheduler.
- Real time threads.
- Real time FIFOs and shared memory.
- Real time interrupt handler.

Yocto

The Yocto Project is an open source project that helps developer create a custom Linux based systems independently from the hardware architecture. [6] Yocto is used to create operating system containing only the features needed by the embedded application.

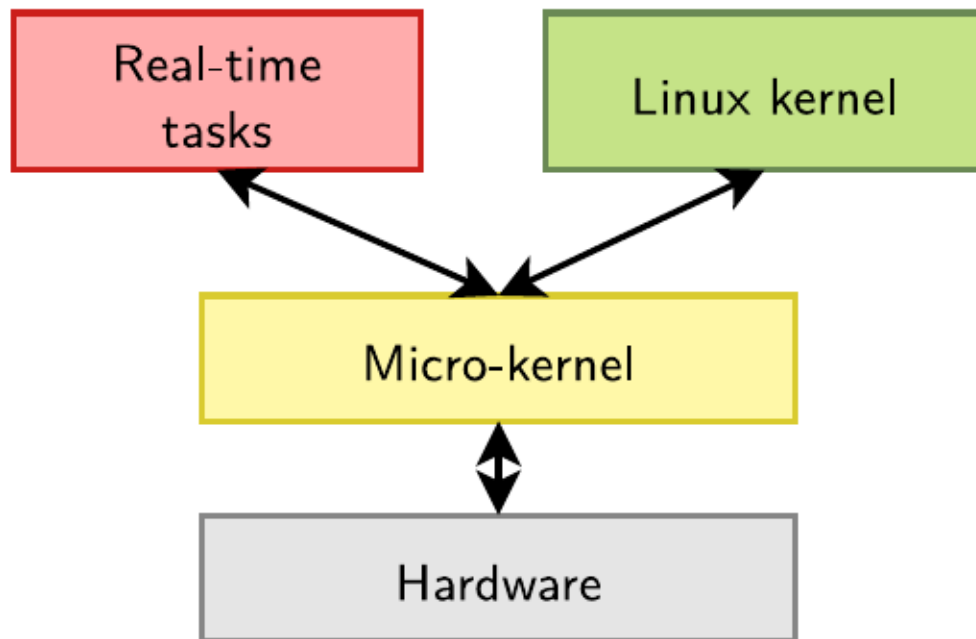


Figure 1.6: Architecture of Real Time Linux

Yocto is characterized by a layer model which grants modularity and the possibility for different developers to work on different part of the system at the same time.

Layers are a set of instructions that tells the build system what to do and they are used to logically separate information in a specific build.

Yocto is formed by two main components: bitbake and OE-Core (Open Embedded core). They are combined to realize the poky host build.

The OpenEmbedded-Core is a collection of information, such as configuration files and recipes used as a common layer to create the custom distributions. It is the starting point from which every distribution is built.

Bitbake is a build engine, which analyze files called recipes, and build an image from them. Inside the recipes there will be all the instruction the engine has to perform to create the image.

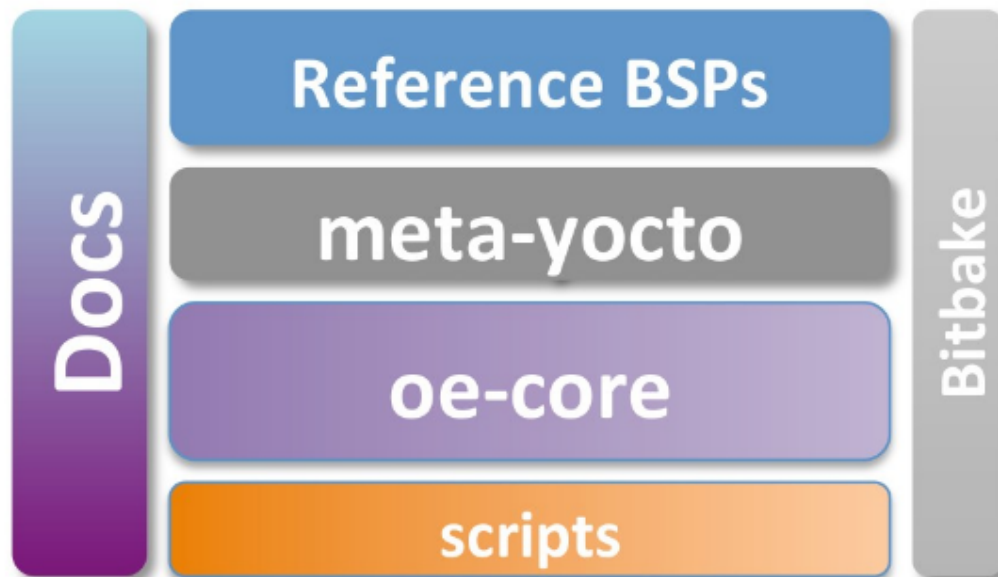


Figure 1.7: Yocto Stack

Bibliography

- [1] Real time linux. <https://www.embedded.com/real-time-linux/>, 2001.
- [2] Embedded linux ubuntu. <https://ubuntu.com/blog/what-is-embedded-linux>, 2019.
- [3] Embedded linux windriver. <https://www.windriver.com/solutions/learning/embedded-linux>, 2020.
- [4] freertos site. <https://www.freertos.org/about-RTOS.html>, 2023.
- [5] Gnu make site. <https://www.gnu.org/software/make/>, 2023.
- [6] Yocto project. <https://www.yoctoproject.org/>, 2023.
- [7] Doug Abbott. *Linux for embedded and real-time applications*. Elsevier, 2011.
- [8] Someone. *Embedded systems architecture*. 2006.